Data colection + Exploratory data analysis

Sergio I. Garcia-Rios

LBJ Survey Design and Analysis

Data Collection + Observational studies and experiments

Use a sample to make inferences about the population

- Ultimate goal: make inferences about populations

- Ultimate goal: make inferences about populations
- Caveat: populations are difficult or impossible to access

- Ultimate goal: make inferences about populations
- Caveat: populations are difficult or impossible to access
- Solution: use a sample from that population, and use statistics from that sample to make inferences about the unknown population parameters

- Ultimate goal: make inferences about populations
- Caveat: populations are difficult or impossible to access
- Solution: use a sample from that population, and use statistics from that sample to make inferences about the unknown population parameters
- The better (more representative) sample we have, the more reliable our estimates and more accurate our inferences will be

- Ultimate goal: make inferences about populations
- Caveat: populations are difficult or impossible to access
- Solution: use a sample from that population, and use statistics from that sample to make inferences about the unknown population parameters
- The better (more representative) sample we have, the more reliable our estimates and more accurate our inferences will be

- Ultimate goal: make inferences about populations
- Caveat: populations are difficult or impossible to access
- Solution: use a sample from that population, and use statistics from that sample to make inferences about the unknown population parameters
- The better (more representative) sample we have, the more reliable our estimates and more accurate our inferences will be

Your Turn

Suppose we want to know how many offspring female squirrels have, on average. It's not feasible to obtain offspring data from on all female squirrels, so we use data from the UT Squirrel Center. We use the sample mean from these data as an estimate for the unknown population mean. Can you see any limitations to using data from the UT Squirrel Center to make inferences about all squirrels?

- When you taste a spoonful of soup and decide the spoonful you tasted isn't salty enough, that's exploratory analysis
- If you generalize and conclude that your entire soup needs salt, that's an inference
- For your inference to be valid, the spoonful you tasted (the sample) needs to be representative of the entire pot (the population)

IDEALLY USE A SIMPLE RANDOM SAMPLE, STRATIFY TO CONTROL FOR A VARIABLE, AND CLUSTER TO MAKE SAMPLING EASIER

Drawing names from a hat

Simple random:

Drawing names from a hat

## Stratified: homogenous strata

Stratify to control for SES

Simple random:
Drawing names from a hat

Stratified: homogenous strata
Stratify to control for SES

Cluster: heterogenous clusters
Sample all chosen clusters


4

Simple random:
Drawing names from a hat

Cluster: heterogenous clusters
Sample all chosen clusters

Stratified: homogenous strata
Stratify to control for SES

Multistage:
Random sample in chosen clusters

4

Your Turn

A city council has requested a household survey be conducted in a suburban area of their city. The area is broken into many distinct and unique neighborhoods, some including large homes, some with only apartments, and others a diverse mixture of housing structures. Which approach would likely be the *least* effective?

(a) Simple random sampling

(b) Stratified sampling, where each stratum is a neighborhood

(c) Cluster sampling, where each cluster is a neighborhood

Your Turn

A city council has requested a household survey be conducted in a suburban area of their city. The area is broken into many distinct and unique neighborhoods, some including large homes, some with only apartments, and others a diverse mixture of housing structures. Which approach would likely be the *least* effective?

(a) Simple random sampling

(b) Stratified sampling, where each stratum is a neighborhood

(c) Cluster sampling, where each cluster is a neighborhood

Sampling schemes can suffer from a variety of biases

- Non-response: If only a small fraction of the randomly sampled people choose to respond to a survey, the sample may no longer be representative of the population

- Non-response: If only a small fraction of the randomly sampled people choose to respond to a survey, the sample may no longer be representative of the population
- Voluntary response: Occurs when the sample consists of people who volunteer to respond because they have strong opinions on the issue since such a sample will also not be representative of the population

- **Non-response:** If only a small fraction of the randomly sampled people choose to respond to a survey, the sample may no longer be representative of the population
- **Voluntary response:** Occurs when the sample consists of people who volunteer to respond because they have strong opinions on the issue since such a sample will also not be representative of the population
- **Convenience sample:** Individuals who are easily accessible are more likely to be included in the sample

A school district is considering whether it will no longer allow high school students to park at school after two recent accidents where students were severely injured. As a first step, they survey parents by mail, asking them whether or not the parents would object to this policy change. Of 6,000 surveys that go out, 1,200 are returned. Of these 1,200 surveys that were completed, 960 agreed with the policy change and 240 disagreed. Which of the following statements are true?

I. Some of the mailings may have never reached the parents.
II. Overall, the school district has strong support from parents to move forward with the policy approval.
III. It is possible that majority of the parents of high school students disagree with the policy change.
IV. The survey results are unlikely to be biased because all parents were mailed a survey.

(a) Only I   (b) I and II   (c) I and III   (d) III and IV   (e) Only IV

A school district is considering whether it will no longer allow high school students to park at school after two recent accidents where students were severely injured. As a first step, they survey parents by mail, asking them whether or not the parents would object to this policy change. Of 6,000 surveys that go out, 1,200 are returned. Of these 1,200 surveys that were completed, 960 agreed with the policy change and 240 disagreed. Which of the following statements are true?

I. Some of the mailings may have never reached the parents.

II. Overall, the school district has strong support from parents to move forward with the policy approval.

III. It is possible that majority of the parents of high school students disagree with the policy change.

IV. The survey results are unlikely to be biased because all parents were mailed a survey.

(a) Only I   (b) I and II   (c) I and III   (d) III and IV   (e) Only IV

Experiments use random assignment to treatment groups, observational studies do not

Facebook Tinkers With Users' Emotions in News Feed Experiment, Stirring Outcry

By VINDU GOEL   JUNE 29, 2014

The New York Times

In an academic paper published in conjunction with two university researchers, the company reported that, for one week in January 2012, it had altered the number of positive and negative posts in the news feeds of 689,003 randomly selected users to see what effect the changes had on the tone of the posts the recipients then wrote.

The researchers found that moods were contagious. The people who saw more positive posts responded by writing more positive posts. Similarly, seeing more negative content prompted the viewers to be more negative in their own posts.

[1]http://www.nytimes.com/2014/06/30/technology/facebook-tinkers-with-users-emotions-in-news-feed-experiment-stirring-outcry.html

Your Turn

A study that surveyed a random sample of otherwise healthy adults found that people are more likely to get muscle cramps when they're stressed. The study also noted that people drink more coffee and sleep less when they're stressed. What type of study is this?

What is the conclusion of the study?

Can this study be used to conclude a causal relationship between increased stress and muscle cramps?

4. Experiments use random assignment to treatment groups, observational studies do not

Your Turn

A study that surveyed a random sample of otherwise healthy adults found that people are more likely to get muscle cramps when they're stressed. The study also noted that people drink more coffee and sleep less when they're stressed. What type of study is this?

Observational

What is the conclusion of the study?

Can this study be used to conclude a causal relationship between increased stress and muscle cramps?

Your Turn

A study that surveyed a random sample of otherwise healthy adults found that people are more likely to get muscle cramps when they're stressed. The study also noted that people drink more coffee and sleep less when they're stressed. What type of study is this?

Observational

What is the conclusion of the study?

There is an association between increased stress & muscle cramps.

Can this study be used to conclude a causal relationship between increased stress and muscle cramps?

Your Turn
A study that surveyed a random sample of otherwise healthy adults found that people are more likely to get muscle cramps when they're stressed. The study also noted that people drink more coffee and sleep less when they're stressed. What type of study is this?

Observational

What is the conclusion of the study?

There is an association between increased stress & muscle cramps.

Can this study be used to conclude a causal relationship between increased stress and muscle cramps?

Muscle cramps might also be due to increased caffeine consumption

Four principles of experimental design: randomize, control, block, replicate

- We would like to design an experiment to investigate if increased stress causes muscle cramps:

- We would like to design an experiment to investigate if increased stress causes muscle cramps:
  - Treatment: increased stress
  - Control: no or baseline stress

- We would like to design an experiment to investigate if increased stress causes muscle cramps:
  - Treatment: increased stress
  - Control: no or baseline stress
- It is suspected that the effect of stress might be different on younger and older people: block for age.

- We would like to design an experiment to investigate if increased stress causes muscle cramps:
    - Treatment: increased stress
    - Control: no or baseline stress
- It is suspected that the effect of stress might be different on younger and older people: block for age.

Why is this important? Can you think of other variables to block for?

Random sampling helps generalizability, random assignment helps causality

	Random assignment	No random assignment	
ideal experiment ↙			↗ *most observational studies*
Random sampling	Causal conclusion, generalized to the whole population.	No causal conclusion, correlation statement generalized to the whole population.	**Generalizability**
No random sampling	Causal conclusion, only for the sample.	No causal conclusion, correlation statement only for the sample.	**No generalizability**
most experiments ↗	Causation	Correlation	↘ *bad observational studies*

1. Use a sample to make inferences about the population
2. Ideally use a simple random sample, stratify to control for a variable, and cluster to make sampling easier
3. Sampling schemes can suffer from a variety of biases
4. Experiments use random assignment to treatment groups, observational studies do not
5. Four principles of experimental design: randomize, control, block, replicate
6. Random sampling helps generalizability, random assignment helps causality

EXPLORATORY DATA ANALYSIS

ALWAYS START YOUR EXPLORATION WITH A VISUALIZATION

Do you see anything out of the ordinary?

How old were you when you had your first kiss?

age at first kiss

Do you see anything out of the ordinary?

How old were you when you had your first kiss?

Some people reported very low ages, which might suggest the survey question wasn't clear: romantic kiss or any kiss?

How are people reporting lower vs. higher values of FB visits?

How many times do you go on Facebook per day?

How are people reporting lower vs. higher values of FB visits?

How many times do you go on Facebook per day?

Finer scale for lower numbers.

Describe the spatial distribution of preferred sweetened carbonated beverage drink.

What is your generic term for a sweetened, carbonated beverage?

- soda
- pop
- coke

Map by Joshua Katz, Department of Statistics, NC State University
Based on survey data from Bert Vaux, Department of Linguistics, University of Cambridge
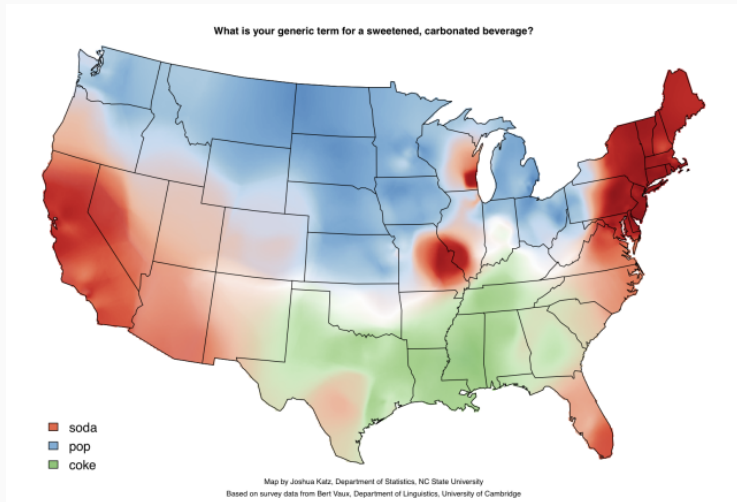
What is missing in this visualization?

What word(s) do you use to address a group of two or more people?

- you guys
- you
- y'all
- you all

Joshua Katz, Department of Statistics, NC State University

When describing numerical distributions discuss shape, center, spread, and unusual observations

- Shape: skewness, modality
- Center: an estimate of a typical observation in the distribution (mean, median, mode, etc.)
    - Notation: μ: population mean, $\bar{x}$: sample mean
- Spread: measure of variability in the distribution (standard deviation, IQR, range, etc.)
- Unusual observations: observations that stand out from the rest of the data that may be suspected outliers

Which of these is most likely to have a roughly symmetric distribution?

(a) salaries of a random sample of people from NY
(b) weights of adult females
(c) scores on an well-designed exam
(d) last digits of phone numbers

Which of these is most likely to have a roughly symmetric distribution?

(a) salaries of a random sample of people from NY
(b) weights of adult females
(c) scores on an well-designed exam
(d) last digits of phone numbers

Your Turn

How do the mean and median of the following two datasets compare?

Dataset 1: 30, 50, 70, 90
Dataset 2: 30, 50, 70, 1000

(a) $\bar{x}_1 = \bar{x}_2$, $median_1 = median_2$
(b) $\bar{x}_1 < \bar{x}_2$, $median_1 = median_2$
(c) $\bar{x}_1 < \bar{x}_2$, $median_1 < median_2$
(d) $\bar{x}_1 > \bar{x}_2$, $median_1 < median_2$
(e) $\bar{x}_1 > \bar{x}_2$, $median_1 = median_2$

Your Turn

How do the mean and median of the following two datasets compare?

Dataset 1: 30, 50, 70, 90
Dataset 2: 30, 50, 70, 1000

(a) $\bar{x}_1 = \bar{x}_2$, $median_1 = median_2$

(b) $\bar{x}_1 < \bar{x}_2$, $median_1 = median_2$

(c) $\bar{x}_1 < \bar{x}_2$, $median_1 < median_2$

(d) $\bar{x}_1 > \bar{x}_2$, $median_1 < median_2$

(e) $\bar{x}_1 > \bar{x}_2$, $median_1 = median_2$

- Most commonly used measure of variability is the standard deviation, which roughly measures the average deviation from the mean
  - Notation: σ: population standard deviation, s: sample standard deviation
- Calculating the standard deviation, for a population (rarely, if ever) and for a sample:

$$\sigma = \sqrt{\frac{\sum_{i=1}^{N}(x_i - \mu)^2}{n}} \qquad s = \sqrt{\frac{\sum_{i=1}^{n}(x_i - \bar{x})^2}{n-1}}$$

- Square of the standard deviation is called the variance.

Why divide by $n - 1$ instead of n when calculating the sample standard deviation?

Why divide by $n-1$ instead of n when calculating the sample standard deviation?

Lose a "degree of freedom" for using an estimate (the sample mean, $\bar{x}$), in estimating the sample variance/standard deviation.

Why divide by $n-1$ instead of n when calculating the sample standard deviation?

Lose a "degree of freedom" for using an estimate (the sample mean, $\bar{x}$), in estimating the sample variance/standard deviation.

Why do we use the squared deviation in the calculation of variance?

Why divide by $n-1$ instead of n when calculating the sample standard deviation?

Lose a "degree of freedom" for using an estimate (the sample mean, $\bar{x}$), in estimating the sample variance/standard deviation.

Why do we use the squared deviation in the calculation of variance?

· To get rid of negatives so that observations equally distant from the mean are weighed equally.
· To weigh larger deviations more heavily.

Our Turn

For the given data set: 7, 6, 5, 5, 9, 10, 11, 10, 9

Calculate

- Range
- Median
- The three quartiles
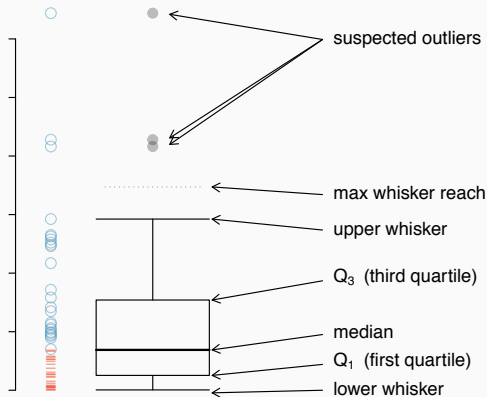- Interquartile range (IQR)
- Draw a boxplot

Robust statistics are not easily affected by outliers and extreme skew

- Mean and standard deviation are easily affected by extreme observations since the value of each data point contributes to their calculation.
- Median and IQR are more robust.
- Therefore we choose median&IQR (over mean&SD) when describing skewed distributions.

Use box plots to display quartiles, median, and outliers

A box plot visualizes the median, the quartiles, and suspected outliers. An outlier is defined as an observation more than $1.5 \times$ IQR away from the quartiles.

suspected outliers

max whisker reach

upper whisker

Q_3 (third quartile)

median

Q_1 (first quartile)

lower whisker

Aplication Exercise

1.1 Distributions of numerical variables

1. Always start your exploration with a visualization
2. When describing numerical distributions discuss shape, center, spread, and unusual observations
3. Robust statistics are not easily affected by outliers and extreme skew
4. Use box plots to display quartiles, median, and outliers